Serial Drivers

by Alessandro Rubini

This article is meant to show the internal structure of device drivers for serial ports, and how they can be perform a variety of services including *ppp* and *slip*. The discussion is based on 2.4 source code, but most of the material applies equally well to 2.2 and 2.0.

The usual view of a serial port

When discussing the software implementation of a serial port, the first thing that comes to mind is `/dev/ttyS0`, as this is the most known character of serial communication, at least on PC-class computers. Since `/dev/ttyS0` is a file special file of type ``char'', a serial driver is often considered a conventional char driver, and phrases like ``char drivers are exemplified by serial ports'' come to mind. Unfortunately, the exemplification is basically wrong.

If you look in the real code, you'll see how the ``char driver'' idea only scratches the surface of what a serial driver is, and that's why a serial driver doesn't quite lend itself to be the prototypical example of char drivers.

Actually, the ``char driver'' abstraction doesn't correctly describe serial device drivers because there is not specific major number associated to each of them. Actually, you can have add-on serial ports that plug in your computer and are managed by specific kernel modules but do not get assigned a new major number.

Looking in `/proc/devices` you'll find that major 4 is associated to the `ttyS` driver, but that's a white lie: text-mode console devices belong to major 4 too, and, actually, Linux-2.0 used the more general ``ttyp'' name for major number 4.

What makes a serial port different from a more conventional char driver like the printer port or a tape drive is its being part of the tty abstraction. Since a serial communication channel can be used to plug an alphanumeric terminal, a serial device driver must be integrated in the terminal emulator layer, called the ``tty'' abstraction, from the name of ancient tele-type devices (still in wide use when Unix was being written)

Overview of tty management

Flexible and powerful tty handling is made up of several building blocks. You must consider that there exist a huge range of tty devices, from VGA and frame-buffer based text consoles up to serial communication channels and virtual terminals as exemplified by (but not limited to) the `xterm` application.

Figure 1 shows the various building blocks that are involved with operation of a serial driver. Most files live in *drivers/char*, if not, the directory specified is relative to the root directory of the kernel source tree. The figure shows how each building block is registered (registration is there in order to allow each block to be implemented as a kernel module).
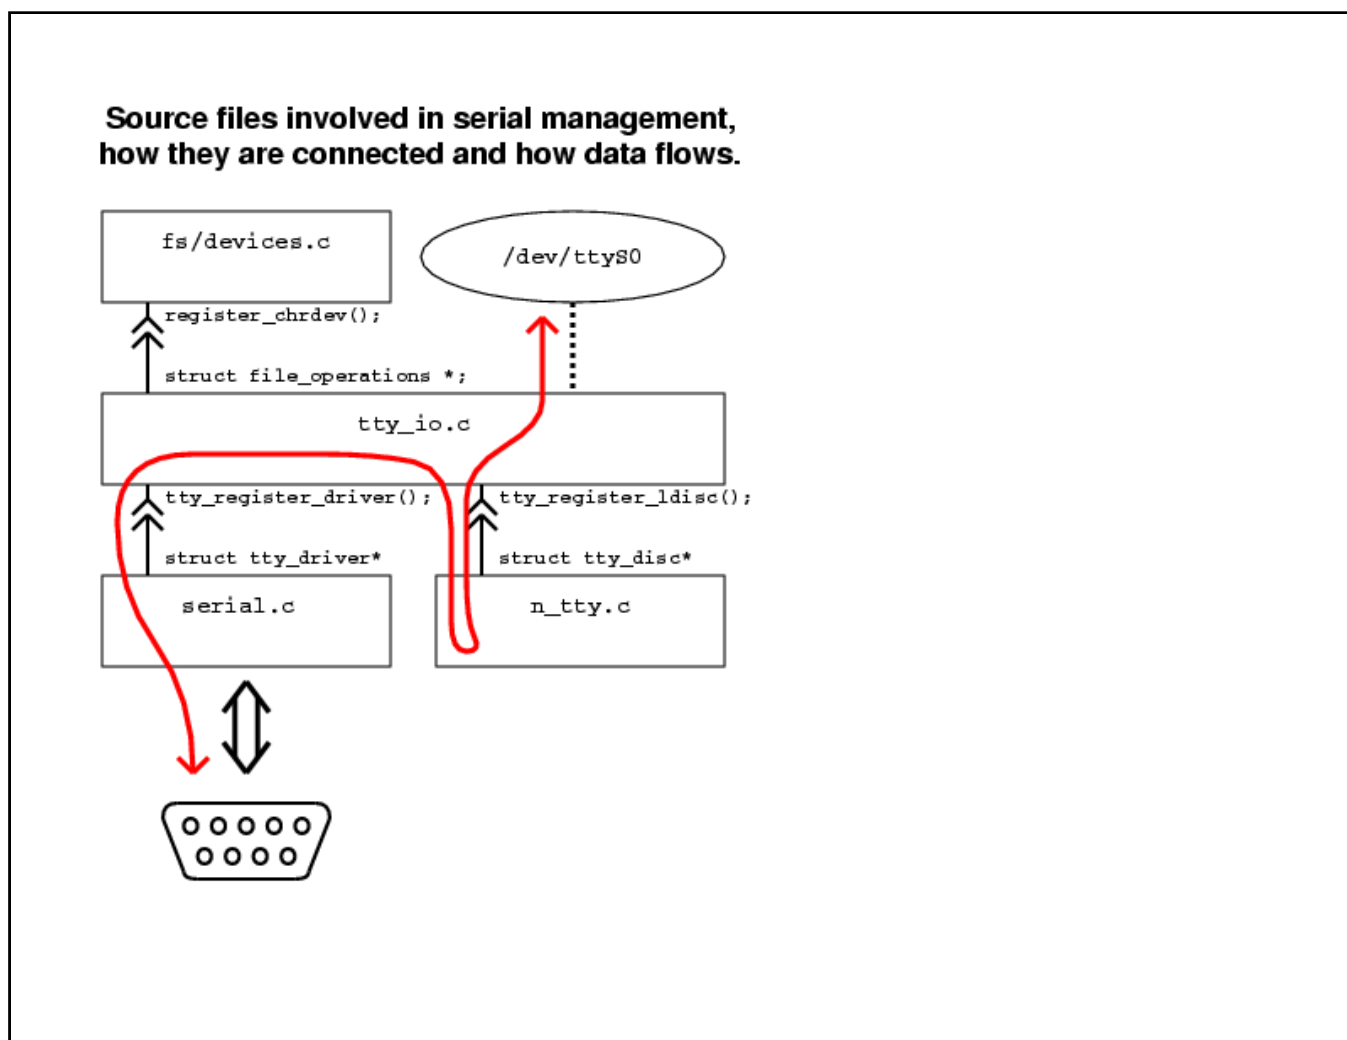

Figure 1

The image is available as PostScript **here**

The file *fs/devices.c* exports the interface used by most system resources to register device drivers, each identified by a major number). This is how

`tty_register_driver` gets hold of a major number if it needs it to support the new tty driver (an object that is introduced below). The function is defined by *tty_io.c*, which also defines the file operations that are used to act on tty devices. File operation are the driver-specific implementations of *read*, *write* and the other system calls that relate to file access.

What *serial.c* or other serial drivers do in order to run their own code is registering a `struct tty_driver` object. The ``driver'' declares the major number and the range of minor numbers it is going to manage, as well as a number of operations it supports. The operations are concerned with input and output of data as well as flow control and communication with higher layers. The implementation of these operations, together with interrupt handling and actual input/output of data is the scope of hardware drivers for serial ports.

The data flow between user space and the serial device driver, therefore, is mediated by the tty layer, that implements functionalities that are common to all tty-type devices. However, not all of tty management is defined be *tty_io.c*, most of the policy is define by the *line discipline*, a software module that disciplines how a physical tty I/O line is used. The default line discipline for Linux is called N_TTY, a name that will be explained later. If *n_tty* is active, input data reaches user space via the usual */dev/* interface and the standard terminal I/O handling (i.e, all the features defined by `<termios.h>`, that make terminal handling so powerful and so difficult).

The red line in figure 1 shows the logical flow of data, from the hardware channel up to the user-accessible device special file and backwards. What keeps it all together is the `struct tty` data item, that in itself includes a pointer to all three relevant objects: the `file_operations` structure that is used in communicating with user space, the `tty_driver` structure that hosts functions to control real hardware, and the `tty_ldisc` structure that lists all entry points to the current line discipline.

Why so complex?

Sure this kind of arrangement may look exceedingly complex. However, as usual, the extra complexity is meant to make things ultimately more flexible and powerful. Adding support for new serial hardware may be different from writing a conventional char driver, but this setup guarantees full tty emulation on all serial ports without any code replication or unneeded complexity in the low-level driver.

Another advantage of this kind of arrangement, possibly even more important than generalized tty support, is in the ability to change the line discipline associated to each tty device. Unlike typical device drivers, whose task is connecting an hardware device with user space, a serial driver has nothing to do with user space; data it receives from the hardware is passed over to the line discipline, and data it receives from software comes via a line discipline method.

A serial driver, therefore, is not concerned in any way with data transfer to/from user space. The task is left to the line discipline, together with all the hairy `termios` handling. This makes it possible for serial data to be steered to a different user-space access facility than its associated *ttyS* device special file.

PPP and slip

Whenever you dial your modem with *ppp* to connect to the Internet, or use the simpler *slip* communication protocol to connect the PC to your Linux palm-top, you are exploiting the complexity just shown. Both *ppp* and *slip* implement their own line discipline; when either of them is run, the tty device is switched to a different line discipline in order to detach */dev/ttyS0* from the serial port and keep all of serial communication within the kernel.

Figure 2 shows the conceptual layout in the *slip* case. I chose not to use *ppp* in the example to avoid extra complexity or incorrect simplification. In version 2.4 of the Linux kernel the *ppp* software implementation is split in several files (once again, it's more structured to be more powerful and avoid code replication across similar devices). The two protocols behave otherwise in the same manner.

Situation when the slip protocol is being used
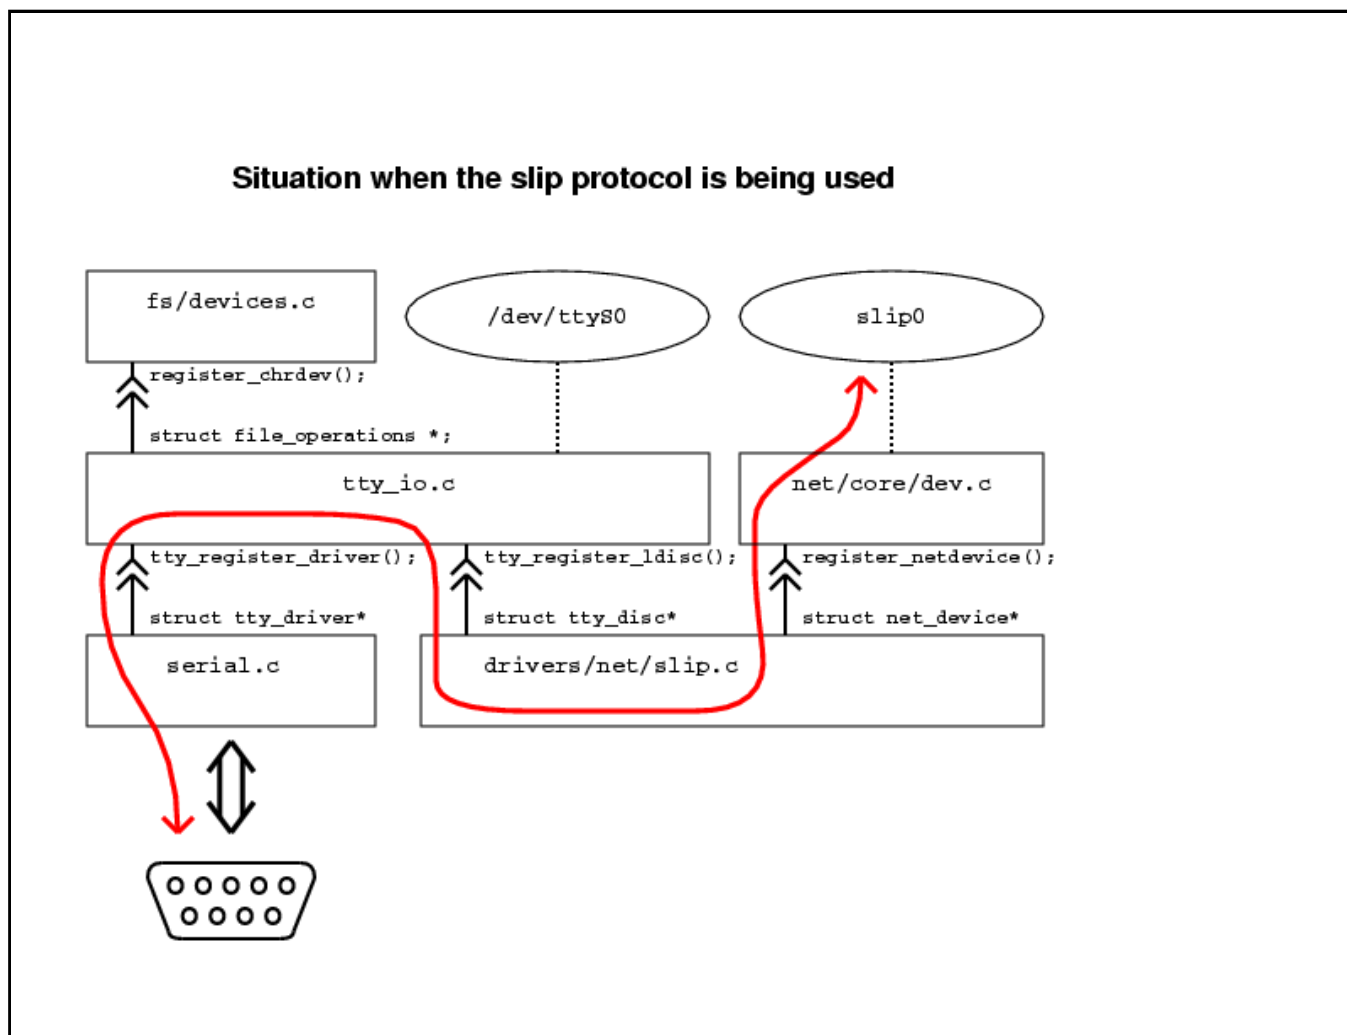
Figure 2

The image is available as PostScript **here**

The role of the *slip* driver, as shown, is registering both a network device (depicted as `slip0` and a line discipline for tty devices (called `N_SLIP`). When the tty device is switched to the new line discipline, TCP/IP communication can begin. The new line discipline sets up data transfer between serial hardware and the network layer; when it is active, nothing can be read from or written to the associated */dev/ttyS* device. As soon as the device (`/dev/ttyS0` or equivalent) is closed, the default line discipline is restored. Actually, that's the main reason why neither *slattach* nor *pppd* can exit after setting up the network channel.

The individual data structures

There are three main data structures involved in tty management (and thus, serial communication):

- `struct tty_struct`: this is the data structure that stays at the core of tty management. It includes both of the following structures. An instance of `tty_struct` is created any time a new tty device is opened, and exists until it is last closed. Note that actual code (in `tty_io.c`) is complicated by the need to preserve `termios` settings across close and open, at least for some of the ttys (like serial ports).
- `struct tty_driver`: this is the low level hardware handling. At open time, the function `get_tty_driver` retrieves the driver for the current tty an places it into the `driver` field of `tty_struct`, where it is further accessed.

```
        struct tty_driver {

            /* the driver states which range of devices it supports */
            short major;          /* major device number */
            short minor_start;    /* start of minor device number*/
            short   num;          /* number of devices */

            /* and has its own operations */
            int  (*open)();
            void (*close)();
            int  (*write)();
            int  (*ioctl)();    /* device-specific control */

            /* return information on buffer state */
            int  (*write_room)();      /* how much can be written */
            int  (*chars_in_buffer)(); /* how much is there to read */

            /* flow control, input and output */
            void (*throttle)();
            void (*unthrottle)();
            void (*stop)();
            void (*start)();

            /* and callbacks for /proc/tty/driver/ */
            int (*read_proc)();
            int (*write_proc)();
        };
```

- `struct tty_ldisc`: the structure is referenced by the `ldisc` field of `tty_struct`. At open time the field is initialized to reference *n_tty*, and user programs can change the current line discipline via *ioctl*, as explained in a while.

```
        struct tty_ldisc {

            /* routines called from above */
            int     (*open)();
            void    (*close)();
            ssize_t (*read)();
            ssize_t (*write)();
            int     (*ioctl)();

            /* routines called from below */
            void    (*receive_buf)();
            int     (*receive_room)();
            void    (*write_wakeup)();
        };
```

The structures are declared in three different files: `tty_struct` is a complex structure defined in *include/linux/tty.h*, a header generally devoted to tty issues. Actually, it is not as interesting to look at as the other two, because user modules rarely need to directly interact with it.

include/linux/tty_driver.h and *include/linux/tty_ldisc.h* are devoted exclusively to the definition of the relevant data structures. The files carry a prominent comment block that explains the exact role of most of the fields. Unlike `tty_struct`, both `tty_driver` and `tty_ldisc` are actively used by authors of device driver modules.

Typically, a kernel module that supports a new kind of hardware transmission will implement a `tty_driver` structure, while a module that uses generic serial hardware for a new purpose will implement a line discipline. For example, if you have a special keyboard that transmits data via a standard RS-232 serial port, you'll need a line discipline that gathers data packets and send them to either the input mechanism (see *drivers/input/input.c* and *include/linux/input.h*) or to the generic keyboard driver (using `handle_scancode()`, exported by *drivers/char/keyboard.c*).

Listing 1 shows the most important operations declared by the `tty_driver` data structure, while listing 2 depicts those exported by the line discipline. Note that those listings are by no means complete, if you look for authoritative information, you should read the relevant header files.
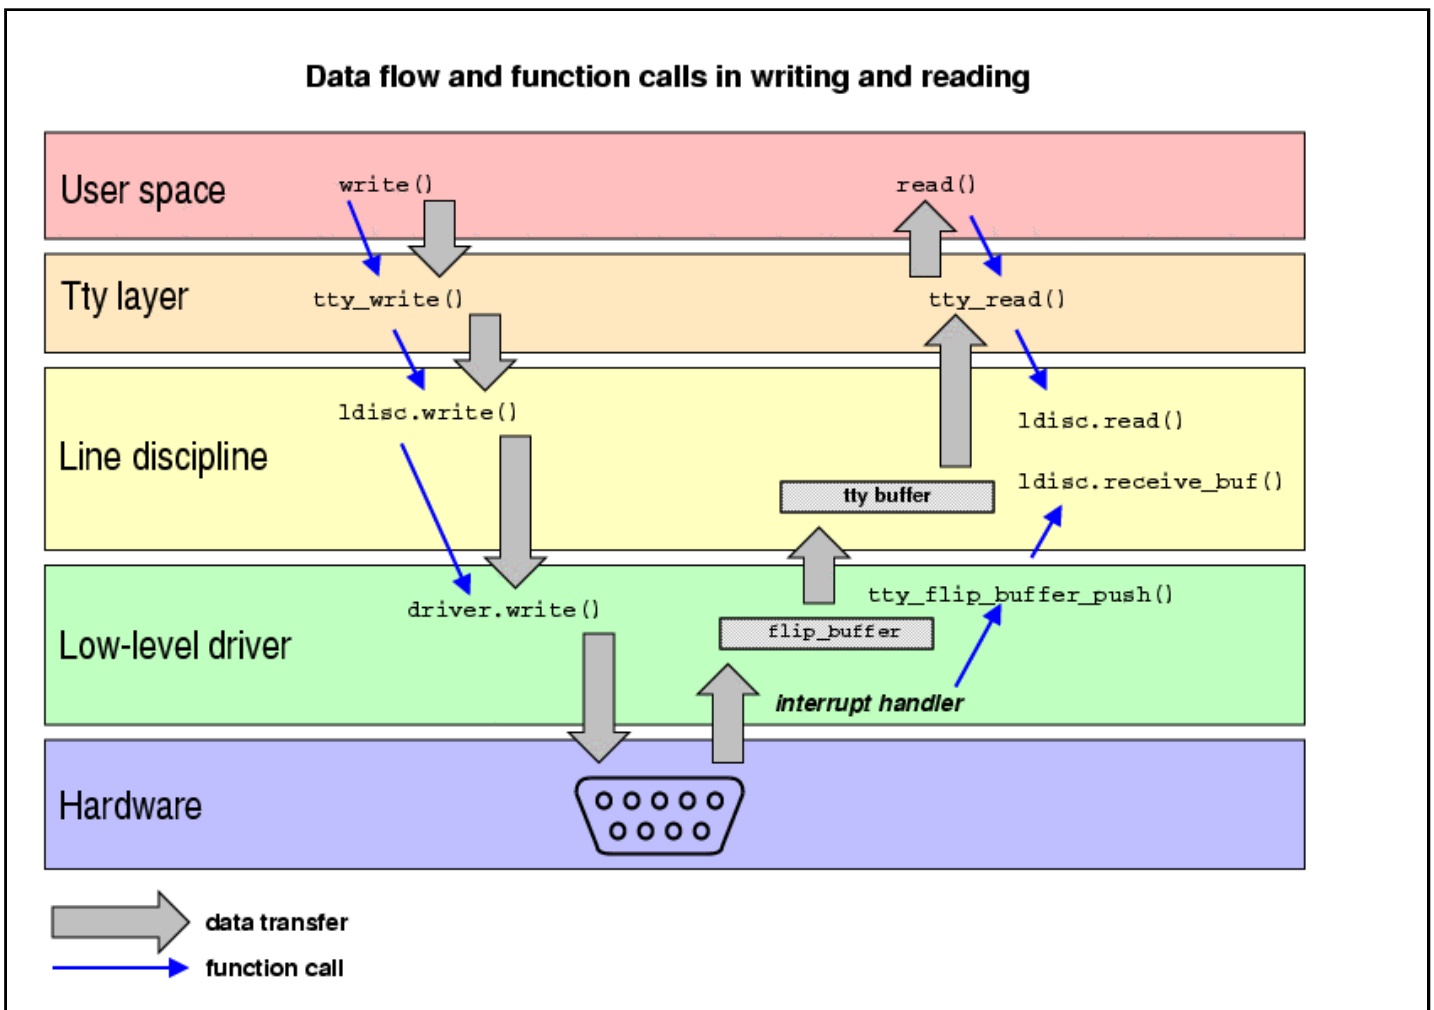
Data flow in reading and writing

Data flow and function calls in writing and reading

Figure 3

The image is available as PostScript **here**

Figure 3 visualizes how data flows from user space down to hardware interfaces and backwards. It refers to the specific case of a standard PC serial port with the default line discipline attached. The logical stacking of line disciplines (near to user space) and tty drivers (near to the hardware itself) should be apparent.

While writing data is straightforward, the reading process may need some explanation. Reading is more complex than writing because there's no direct causal relationship between hardware (that pushes up data when it arrives) and user space (that requests data when it needs it). The solution is, as you expected, use of buffering: data received by hardware remains on hold in a kernel buffer until a user space program requests it; whenever a user program asks for data and the buffer is empty, the program is put to sleep, and is awaken only when the buffer is filled with at least partially. Note that a write buffer exists as well, however, the *write* implementation is much more straightforward because each step towards hardware level is directly driven by the step above it. There are no uncontrolled delays in data transmission, and the buffer is only needed to decouple hardware transmission from program flow. The figure does not show it for simplicity.

When tty devices are concerned, the read buffer lives within the tty data structure; while this makes `struct tty_struct` considerably bigger, there is no point in keeping the buffer elsewhere: each tty can't transfer data without a buffer, and tty devices are dynamically allocated so no memory is wasted in buffers for unused devices.

Actually, tty-related buffering is organized in two levels: kernel developers chose to provide both a ``conventional'' buffer, where data is waiting to be eaten by the line discipline (i.e., in the default case, being transferred to user space), and a ``flip'' buffer, used by hardware routines to store incoming data as quickly as possible, without the need to synchronize for concurrent access: flip buffers are exclusive ownership of the hardware device, which eventually calls *tty_flip_buffer_push* to deliver data to the tty buffer, where the line discipline pulls it from.

It's interesting to note that the flip buffer is laid out as two physical buffers that are alternatively written to. This allows more reliable operation, as the interrupt handler will always have a whole buffer available for writing. The function *flush_to_ldisc*, called by the low-level driver and part of the tty layer (i.e., *tty_io.c*), arranges for the flip buffer to be flipped, before the interrupt handler returns. This layout, by the way, is why the flip buffer is called so.

How to use a custom line discipline

Kernel code, as stated, can register a new line discipline with the tty subsystem, and this is also available to modularized code. You could, therefore, write your own line discipline and register it. Each line discipline is identified by a number, and a symbolic name for it is available, as common with C programming. Assigned numbers are given a name by *include/asm/termios.h*.

The default tty line discipline is identified by a number of `N_TTY`, PPP uses `N_PPP` and so on. Unfortunately, no line discipline numbers have currently been reserved for ``local use'', so you can only experiment with the numbers that are not used on your system. Actually, no official driver currently used `N_MOUSE`, so this is a good bet for your custom line discipline.

In order to activate the `N_MOUSE` line discipline, the user space program must use this code:

```
#include

int i = N_MOUSE;
ioctl(fd, TIOCSETD, &i);
```

The role of register_serial

If you noted that *drivers/char/serial.c* exports a function called *register_serial*, you may wonder what's its role in the tty architecture just outlined. As a matter of facts, the facility is only an hook offered by the ``standard'' serial tty driver in order to easy run-time addition of standard serial ports. The ``serial'' being registered is not a whole software module, but rather, only a definition or parameters to use for the new serial port. The parameters are described by `struct serial_struct`, which in turn is defined by *include/linux/serial.h*; they are used by the conventional serial driver, exploiting the de-facto standardization that exists on PC serial ports. You can't use the function to register a driver for serial hardware of a different kind than a 16450 or compatible UART. The list of supported hardware for the PC platform is found in the array `uart_config`, in *serial.c*. Other platforms offer different implementations for `register_serial`.

The serial console

Serial ports can also be used as console devices, and this kind of functionality is separate from tty management. Actually, the console device sits in the lowest levels of Linux, as it must bring critical information out of the system as soon as possible. It just can't be involved in all the complexity of tty management. But in version 2.4 (as well as 2.2) console management is a world of its own, with its own data structures and functions that make it very flexible yet reliable. Discussion of such world is best delayed to next month's column.

Alessandro is an independent consultant based in Italy. Even though he claims to be a programmer, his activity consists mostly in advocating libre software and documenting what real programmers do. He runs the GNU operating system on his computers, and can thus be reached as `rubini@gnu.org`.